

```
In [78]: import IPython.display as display
import PIL import Image
import matplotlib.pyplot as plt
import numpy as np
import os
import pathlib
import scipy.io
import tensorflow as tf

from tensorflow.keras.layers import Dense, Conv2D, Flatten, Dropout, Max
Pooling2D, AveragePooling2D
```

Load Data

```
In [26]: NUM_TRAIN_SAMPLES = 73257
NUM_TEST_SAMPLES = 26032
IMAGE_SIZE = 32
RGB = 3
```

```
In [27]: train = scipy.io.loadmat('data/mat/train_32x32.mat')
X_train = np.asarray([train['X'][:, :, :, i] for i in range(NUM_TRAIN_SAMPLES)])
y_train = train['y']
```

```
In [28]: print ('X train: {0} y train: {1}'.format(X_train.shape, y_train.shape))

X train: (73257, 32, 32, 3) y train: (73257, 1)
```

```
In [29]: test = scipy.io.loadmat('data/mat/test_32x32.mat')
X_test = np.asarray([test['X'][:, :, :, i] for i in range(NUM_TEST_SAMPLES)])
y_test = test['y']
```

```
In [30]: print ('X test: {0} y test: {1}'.format(X_test.shape, y_test.shape))

X test: (26032, 32, 32, 3) y test: (26032, 1)
```

```
In [31]: # Normalize pixel values to be between 0 and 1
X_train, X_test = X_train / 255.0, X_test / 255.0
```

```
In [32]: plt.figure(figsize=(10,10))  
for i in range(25):  
    plt.subplot(5,5,i+1)  
    plt.xticks([])  
    plt.yticks([])  
    plt.grid(False)  
    plt.imshow(X_train[i], cmap=plt.cm.binary)  
plt.show()
```



Define Architecture

3-Layer CNN with 3x3 filter size and 32, 64, and 128 filters.

Alternating 2x2 average pooling layers between each convolution

ReLu activation

One softmax and output layer

```
In [92]: model = tf.keras.Sequential([
            Conv2D(32, (3, 3), activation='relu', input_shape=(IMAGE_SIZE, I
            IMAGE_SIZE, 3)),
            AveragePooling2D((2, 2)),
            Conv2D(64, (3, 3), activation='relu'),
            AveragePooling2D((2, 2)),
            Conv2D(128, (3, 3), activation='relu'),
            AveragePooling2D((2, 2))
        ])
```

```
In [93]: model.summary()
```

Model: "sequential_23"

Layer (type)	Output Shape	Param #
=====		
conv2d_139 (Conv2D)	(None, 30, 30, 32)	896
average_pooling2d_6 (Average)	(None, 15, 15, 32)	0
conv2d_140 (Conv2D)	(None, 13, 13, 64)	18496
average_pooling2d_7 (Average)	(None, 6, 6, 64)	0
conv2d_141 (Conv2D)	(None, 4, 4, 128)	73856
average_pooling2d_8 (Average)	(None, 2, 2, 128)	0
=====		
Total params: 93,248		
Trainable params: 93,248		
Non-trainable params: 0		

```
In [94]: model.add(Flatten())
          model.add(Dense(64, activation='relu'))
          model.add(Dense(11))
```

```
In [95]: model.summary()
```

```
Model: "sequential_23"
```

Layer (type)	Output Shape	Param #
=====		
conv2d_139 (Conv2D)	(None, 30, 30, 32)	896
average_pooling2d_6 (Average)	(None, 15, 15, 32)	0
conv2d_140 (Conv2D)	(None, 13, 13, 64)	18496
average_pooling2d_7 (Average)	(None, 6, 6, 64)	0
conv2d_141 (Conv2D)	(None, 4, 4, 128)	73856
average_pooling2d_8 (Average)	(None, 2, 2, 128)	0
flatten_6 (Flatten)	(None, 512)	0
dense_12 (Dense)	(None, 64)	32832
dense_13 (Dense)	(None, 11)	715
=====		
Total params: 126,795		
Trainable params: 126,795		
Non-trainable params: 0		

```
In [96]: model.compile(optimizer='adam',
                        loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                        metrics=['accuracy'])
```

Train Model for 10 epochs

```
In [97]: history = model.fit(X_train, y_train, epochs=10,  
                             validation_data=(X_test, y_test))
```

Train on 73257 samples, validate on 26032 samples

Epoch 1/10

73257/73257 [=====] - 28s 383us/sample - loss: 1.1406 - accuracy: 0.6199 - val_loss: 0.5682 - val_accuracy: 0.8356

Epoch 2/10

73257/73257 [=====] - 29s 395us/sample - loss: 0.4902 - accuracy: 0.8549 - val_loss: 0.4638 - val_accuracy: 0.8661

Epoch 3/10

73257/73257 [=====] - 29s 393us/sample - loss: 0.4001 - accuracy: 0.8818 - val_loss: 0.4176 - val_accuracy: 0.8789

Epoch 4/10

73257/73257 [=====] - 28s 388us/sample - loss: 0.3494 - accuracy: 0.8971 - val_loss: 0.3764 - val_accuracy: 0.8906

Epoch 5/10

73257/73257 [=====] - 27s 373us/sample - loss: 0.3159 - accuracy: 0.9067 - val_loss: 0.3677 - val_accuracy: 0.8967

Epoch 6/10

73257/73257 [=====] - 27s 373us/sample - loss: 0.2923 - accuracy: 0.9135 - val_loss: 0.3365 - val_accuracy: 0.9051

Epoch 7/10

73257/73257 [=====] - 27s 373us/sample - loss: 0.2726 - accuracy: 0.9196 - val_loss: 0.3237 - val_accuracy: 0.9095

Epoch 8/10

73257/73257 [=====] - 27s 371us/sample - loss: 0.2571 - accuracy: 0.9239 - val_loss: 0.3309 - val_accuracy: 0.9072

Epoch 9/10

73257/73257 [=====] - 27s 370us/sample - loss: 0.2414 - accuracy: 0.9292 - val_loss: 0.3179 - val_accuracy: 0.9115

Epoch 10/10


73257/73257 [=====] - 28s 379us/sample - loss: 0.2279 - accuracy: 0.9332 - val_loss: 0.3280 - val_accuracy: 0.9091

Evaluate Model - 90% accuracy

```
In [98]: plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.5, 1])
plt.legend(loc='lower right')
```

```
test_loss, test_acc = model.evaluate(X_test, y_test, verbose=2)
```

```
26032/26032 - 3s - loss: 0.3280 - accuracy: 0.9091
```



```
In [99]: test_acc
```

```
Out[99]: 0.90911186
```

Review

1. Normalized images during pre-processing
2. Experimented with varying number of layers and found that a three-layer convolutional network alternating pooling layers provided the best results
3. Experimented with 32, 64, and 128 filter sizes to extract more detail from each layer
4. Applied a 3x3 filter size for each convolutional layer
5. Used an average pooling layer between each convolution because contrast between house numbers and backgrounds were not clearly defined
6. Used ReLu activation function to quickly train the model
7. Achieved 90% accuracy

```
In [ ]:
```